

Introduction

Since I [already programmed](#) the IBM Model 1 for a course on Statistical Machine Translation, the submitted code is just a polished and commented version of it. I am also attaching a draft of the paper *Leveraging Neural Machine Translation for Word Alignment*, which you may find interesting. Please do not share it outside of your team as it is currently in a review process. It covers three topics: (1) using an NMT to get word alignments, (2) various hard alignment extractor methods and (3) combination of all in a small neural network. The extractor methods are in brackets. Hard alignment extractors are reused from this projects and this document describes them very superficially and without mathematical notation. The original version was based on Kohen's pseudocode, which he recently [revised](#). For the graphics I also used [Slow Align](#), a small tool for displaying word alignment (needs major reworking).

The submitted files:

```
img/
  align_{a3,a4,int}.svg # illustrations for this document
src/
  aligner.py             # main aligner code
  run_eval.py            # runs a configuration and evaluates it
  fastalign.sh           # runs and evaluates fast_align configurations
  extractor.py           # hard alignment extractors, re-used from LeverageAlign
  intersector.py         # used for intersection alignments extrensicallly
  intersect_reverse.py   # computes forward and backward alignments and intersects them
Leverage_Align.pdf       # draft of Leveraging Neural Machine Translation for Word Alignment paper
rust/main.rs            # rust implementation of IBM1
```

Baseline

Baseline observed.

```
Precision = 0.243110
Recall     = 0.544379
AER        = 0.681684
```

Seventh sentence from the dataset. It has complete disregard for the position of the tokens and also quite low precision.

Alignment 7 KEY: () = guessed, * = sure, ? = possible

[illegible]

Custom Models

IBM1 (A0)

Even though the code uses numpy whenever possible, it still takes considerable amount of time to run on all 100k examples (~10 minutes for 5 steps). This version simply takes the maximum aligned token. This makes a strong assumption that every token is aligned to exactly one other token. Lowercasing the whole corpus led to slightly better results, therefore I report performance on lowercased input.

Results for IBM1 / A0 with 5 EM steps.

```
python3 ./src/aligner.py -n 100000 -e A0 -s 5 | python3 ./jhu-mt-hw/hw2/score-alignments --data jhu-mt-hw/hw
Precision = 0.627837
Recall    = 0.786982
AER       = 0.318318
```

IBM1 + NULL (A0+NULL)

I also experimented with adding NULL. It was added to the target sentence and then went through the rest of the pipeline. Before output all alignments to this token were cancelled. The results were, however, worse than without it.

```
Precision = 0.618380
Recall    = 0.742604
AER       = 0.338776
```


Threshold (A2)

Results for A2(0.35) with 7 EM steps.

```
python3 ./src/aligner.py -n 100000 -e A2 -s 7 | python3 ./jhu-mt-hw/hw2/score-alignments --data jhu-mt-hw/hw
Precision = 0.685567
Recall    = 0.778107
AER       = 0.280435
```

Proportion of best (A3)

This extraction method examines all target tokens and selects the best one. All other alignments which have scores of at least alpha times the maximum one are taken. Here alpha is a parameter. In the following example consider alpha=0.4, then *menu* would align to *Popupmenü* and *Dialogfeld* (mistakenly) and not to others.




Results for A3(0.98) with 7 EM steps.

```
python3 ./src/aligner.py -n 100000 -e A3 -s 7 | python3 ./jhu-mt-hw/hw2/score-alignments --data jhu-mt-hw/hw
Precision = 0.538631
Recall    = 0.878698
AER       = 0.368971
```

Proportion of reverse best (A4)

The same approach can be done from the target token perspective. Surprisingly, this works a bit better than A3. In the example for alpha=0.5, *Popupmenü* would align to *pop-up*, *menu* and *dialog* (mistakenly).



Results for A4(0.98) with 7 EM steps.

Intersection (A2*A3*A4)

Choose Export from the pop-up menu at the bottom of the dialog box .

Wählen Sie im Pop-upmenü unten im Dialogfeld die Option " Exportieren . "

```
python3 ./src/aligner.py -n 100000 -e A2*A3*A4 -s 7 | python3 ./jhu-mt-hw/hw2/score-alignments --data jhu-mt
Precision = 0.758483
Recall    = 0.784024
AER       = 0.231228
```

The final step combines the previous approach by computing the alignment from the other direction and intersecting it as well. In this version only alignments which are within 7 places of each other are considered. This promotes diagonal alignment. A more elegant solution would be to not create this hard boundary, but only increase scores to alignments on the diagonal. The decision process would then be delegated to the extractors.

```
src/intersect_reverse.sh
Precision = 0.879781
Recall    = 0.718935
AER       = 0.197443
```

Alignment 7 KEY: () = guessed, * = sure, ? = possible

(*)																					je
*	?	(?) ?																			le
?	?																				ai
																					cité
																					comme
						(*)															exemple
						?	(*)														de
																					propriété
																					publique
																					créatrice
																					,
																					dynamique
																					et
																					efficace
																					.
I	h	s	t	C	a	a	e	o	c	,	a	e	p	o	.						
a	v	r	h	D	s	n	x	f	r		g	f	u	w							
v	e	e	e	C			a		e		g	f	b	n							
		s					m		a		r	e	l	e							
		s						t		e	c	i	r								
		s							s	t	c	s									
		d							s	i		h									
										v	e		i								
											e			p							

fast_align

The script `src/fastalign.sh` runs `fast_align` with best practice arguments (`-v` Dirichlet prior and `-do` close to diagonal). EM steps are default 5.

```
Precision = 0.788122
Recall     = 0.831361
AER        = 0.196670
```

Alignment of the seventh sentence in the data by fast_align. It performed slightly better overall (by 0.001). I am however sure, that if one would choose better parameters, then both fast_align and the proposed aligner would produce even better results.

Alignment 7 KEY: () = guessed, * = sure, ? = possible

(*)																				je
	(?) ?																			le
(*) ?																				ai
? (?)	()																			cité
		(*)																		comme
			(?)(*)																	exemple
				(*)																de
																				propriété
																				publique
								(*)	()											créatrice
									(*)											,
										(*)										dynamique
																				et
																				efficace
																				.
I	h	s	t	C	a	a	e	o	c	,	a	e	p	o	.					
a	v	r	e	D	s	n	x	f	r		g	f	u	w						
e		e		C			a		e		r	b	n							
		s					m		a		e	i	r							
		s					e		t		s	c	s							
		e							i		s	t	h							
		d							v		i	v	i							
									e		e		p							

Without -d the results are:

```
Precision = 0.645161
Recall    = 0.766272
AER       = 0.313448
```

Performance

I was quite concerned with the implementation speed. So I wrote it in Rust (rust/). The time went from ~10min to ~10.2s. This is better than fast_align (~12.9s). The comparison is however not fair from both sides, since the Rust implementation uses only IBM1 EM computation, but on the other multi-threads only the maximization step. Also these measurements should be taken with a bag of salt, since it is only two hot passes on a development notebook.

Only the A0 extraction method was implemented. The results are almost the same as for the IBM1 / A0 model. The small differences are caused by unstable argmax.

```
Precision = 0.618585
Recall     = 0.804734
AER       = 0.322002
```